



# **Title:** **Vision AI in 5 Days:** **From Beginner to** **Image Recognition** **Expert**

Subtitle:

*Image Classification using CNN in Python*

*Presented by:*

*Gayatri Sharma*

# Dataset: CIFAR-10 / Cats vs Dogs

Total images:(e.g., 50,000 train, 10,000 test)

Processing Steps :

Image resizing

Normalization

Train-Test split

Data augmentation  
(rotation, flip, zoom)

# Dataset Preview



```
import tensorflow as tf
from tensorflow.keras import layers, models
import matplotlib.pyplot as plt
import numpy as np

(x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0 # Normalize pixel values

class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
               'dog', 'frog', 'horse', 'ship', 'truck']

model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),

    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Conv2D(64, (3, 3), activation='relu'),
```



```
layers.MaxPooling2D((2, 2)),

layers.Conv2D(64, (3, 3), activation='relu'),
layers.MaxPooling2D((2, 2)),

layers.Conv2D(64, (3, 3), activation='relu'),

layers.Flatten(),
layers.Dense(64, activation='relu'),
layers.Dense(10)
])

model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])

history = model.fit(x_train, y_train, epochs=10,
                    validation_data=(x_test, y_test), batch_size=64)

test_loss, test_acc = model.evaluate(x_test, y_test)
print(f"\n✅ Test Accuracy: {test_acc * 100:.2f}%")
print(f"❌ Test Loss: {test_loss:.4f}")
```



Activate Windows

Go to Settings to activate Windows.



Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>



170498071/170498071 2s 0us/step

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base\_conv.py:113: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer that does not require it.  
super().\_\_init\_\_(activity\_regularizer=activity\_regularizer, \*\*kwargs)

Epoch 1/10

782/782 62s 77ms/step - accuracy: 0.3299 - loss: 1.8096 - val\_accuracy: 0.5238 - val\_loss: 1.3540

Epoch 2/10

782/782 82s 78ms/step - accuracy: 0.5534 - loss: 1.2538 - val\_accuracy: 0.6114 - val\_loss: 1.0989

Epoch 3/10

782/782 80s 75ms/step - accuracy: 0.6138 - loss: 1.0870 - val\_accuracy: 0.6368 - val\_loss: 1.0201

Epoch 4/10

782/782 59s 75ms/step - accuracy: 0.6599 - loss: 0.9610 - val\_accuracy: 0.6518 - val\_loss: 0.9905

Epoch 5/10

782/782 59s 75ms/step - accuracy: 0.6861 - loss: 0.8968 - val\_accuracy: 0.6723 - val\_loss: 0.9581

Epoch 6/10

782/782 81s 73ms/step - accuracy: 0.7118 - loss: 0.8233 - val\_accuracy: 0.6787 - val\_loss: 0.9379

Epoch 7/10

782/782 60s 77ms/step - accuracy: 0.7250 - loss: 0.7764 - val\_accuracy: 0.6921 - val\_loss: 0.8960

Epoch 8/10

782/782 61s 78ms/step - accuracy: 0.7493 - loss: 0.7187 - val\_accuracy: 0.6976 - val\_loss: 0.8721

Epoch 9/10

782/782 58s 74ms/step - accuracy: 0.7611 - loss: 0.6819 - val\_accuracy: 0.6922 - val\_loss: 0.8925

Epoch 10/10

782/782 60s 77ms/step - accuracy: 0.7741 - loss: 0.6457 - val\_accuracy: 0.7126 - val\_loss: 0.8589

313/313 3s 11ms/step - accuracy: 0.7187 - loss: 0.8484

✓ Test Accuracy: 71.26%

✗ Test Loss: 0.8589

Activate Windows

Go to Settings to activate Windows.

# Model Architecture:

- **Model Type:** Convolutional Neural Network (CNN)
- **Layers:** Conv → Pool → Conv → Pool → Dense → Output
- **Activation Functions:** ReLU, Softmax
- **Optimizer:** Adam
- **Loss Function:** Categorical Crossentropy

# Results & Evaluation

- **Training Accuracy:** ~77%
- **Validation Accuracy:** ~71%
- **Test Accuracy:** 71.26%
- **Test Loss:** 0.8589
- **Visuals:**
  - Accuracy & Loss curves (Matplotlib se plot)
  - Confusion matrix (agar banayi ho)
- Optional: Precision, Recall, F1-score table.

# • Conclusion & Future Scope

## **Conclusion:**

- CNN model successfully trained for image classification.
  - Achieved 71% accuracy on test dataset.
  - **Future Improvements:**
- Use Transfer Learning (MobileNetV2, ResNet)
  - More data augmentation
  - Hyperparameter tuning



# Demo Video :

- <https://drive.google.com/file/d/1SyWFwV90ObJnLJc9OjXRXUzF151w24y4/view?usp=sharing>